

NATURAL LANGUAGE PROCESSING AND TEXT CLASSIFICATION

Juan Ginzo

NATURAL LANGUAGE PROCESSING AND TEXT CLASSIFICATION

LEARNING OBJECTIVES

- Define natural language processing
- List common tasks associated with
 - use-cases
 - tokenization
 - tagging
 - parsing
- Demonstrate how to classify text or documents using scikit-learn

COURSE

PRE-WORK

PRE-WORK REVIEW

- Experience with scikit-learn classifiers, specifically random forests and decision trees
- Install the Python package spacy with `pip install spacy`
- Run the spacy download data command

```
python -m spacy.en.download --force all
```

OPENING


NATURAL LANGUAGE PROCESSING AND TEXT CLASSIFICATION

REVIEW: DECISION TREES AND RANDOM FORESTS

- What are decision trees?
- What are random forests?

REVIEW: DECISION TREES AND RANDOM FORESTS

- Decision trees are models that ask a series of questions. The next question depends upon the answer to the previous question.
- Random forest models are ensembles of decision trees that are randomized in the way they are created.



REVIEW: DECISION TREES AND RANDOM FORESTS

- Decision trees are weak learners that are easy to overfit.
- Random forests are strong models that are made up of a collection of decision trees.
 - They are non-linear (as opposed to logistic regression).
 - They are mostly black-boxes (no coefficients, although we do have a measure of feature importance).
 - They can be used for classification or regression.

INTRODUCTION

NATURAL LANGUAGE PROCESSING

WHAT IS NATURAL LANGUAGE PROCESSING (NLP)?

- Natural language processing is the task of extracting meaning and information from text documents.
- There are many types of information we might want to extract.
- These tasks may range from simple classification tasks, such as deciding what category a piece of text falls into, to more complex tasks like translating or summarizing text.
- For most tasks, a fair amount of pre-processing is required to make the text digestible for our algorithms. We typically need to *add structure* to our *unstructured data*.

WHAT IS NATURAL LANGUAGE PROCESSING (NLP)?

- Many AI assistant systems are typically powered by fairly advanced NLP engines.
- A system like Siri uses voice-to-transcription to record a command and then various NLP algorithms to identify the question asked and possible answers.



TASKS IN NLP

- In the context of creating predictive models, one of our main tasks is to be able to extract features from the text.
- We need to:
 - Extract meaningful features (words) from the text (tokenization).
 - Normalize those words to a canonical form. (stemming/lemmatization)

TOKENIZATION

- Tokenization is the task of separating a sentence into its constituent parts, or *tokens*.
- Determining the “words” of a sentence seems easy but can quickly become complicated with unusual punctuation (common in social media) or different language conventions.

TOKENIZATION

- What sort of difficulties can you find in the following sentence?
- The L.A. Lakers won the NBA championship in 2010, defeating the Boston Celtics.

TOKENIZATION EXAMPLES

My house is located in Uptown. → [My, house, is, located, in, Uptown]

The Lakers are my favorite team. → [The, Lakers, are, my, favorite, team]

Data Science is the future! → [Data, Science, is, the, future]

GA has many locations. → [GA, has, many, locations.]

LEMMATIZATION AND STEMMING

- How would you describe the relationship between the terms ‘bad’ and ‘badly’ or ‘different’ and ‘differences’?
- *Stemming* and *lemmatization* help identify common roots of words.
- *Stemming* is a crude process of removing common endings from sentences, such as ‘s’, ‘es’, ‘ly’, ‘ing’, and ‘ed’.

LEMMATIZATION AND STEMMING

- Lemmatization is a more refined process that uses specific language and grammar rules to derive the root of a word.
- This is useful for words that do not share an obvious root such as ‘better’ and ‘best’.
- What are some other examples of words that do not share an obvious root?

LEMMATIZATION AND STEMMING EXAMPLES

Lemmatization

shouted → shout

best → good

better → good

good → good

wiping → wipe

hidden → hide

Stemming

badly → bad

computing → comput

computed → comput

wipes → wip

wiped → wip

wiping → wip

ACTIVITY: KNOWLEDGE CHECK



EXERCISE

ANSWER THE FOLLOWING QUESTIONS


1. What other words or phrases might cause problems with stemming? Why?
2. What other words or phrases might cause problems with lemmatization? Why?

DELIVERABLE

Answers to the above questions

PARSING AND TAGGING

- ▶ Another important task in NLP is parsing and tagging.
- ▶ In order to understand the various elements of a sentence, we need to *tag* important topics and *parse* their dependencies.
- ▶ Our goal is to identify the *actors* and *actions* in the text in order to make informed decisions.



PARSING AND TAGGING

- If we are processing financial news, we might need to identify which companies are involved and which actions they are taking.
- If we are writing an assistant application, we might need to identify specific command phrases in order to determine what is being asked:
 - e.g. “Siri, when is my next appointment?”

PARSING AND TAGGING

- Tagging and parsing is made up of a few overlapping subproblems:
 - **“Parts of speech” tagging**: What are the parts of speech in a sentence (e.g. noun, verb, adjective, etc)?
 - **Chunking**: Can we identify the pieces of the sentence that go together in meaningful chunks (e.g. noun or verb phrases)?
 - **Named entity recognition**: Can we identify *specific* proper nouns? Can we pick out people and locations?

ACTIVITY: KNOWLEDGE CHECK



EXERCISE

ANSWER THE FOLLOWING QUESTIONS

1. How might NLP be applied within your current jobs or final projects?
2. What are some other potential NLP use-cases?

DELIVERABLE

Answers to the above questions

DEMO

NATURAL LANGUAGE PROCESSING WITH 'SPACY'

NATURAL LANGUAGE PROCESSING WITH 'SPACY'

- Most NLP techniques require pre-processing large collections of annotated text in order to learn specific language rules.
- There are many tools available for English and other popular languages.
- Each tool typically requires a large amount of data and large databases of special use-cases, including language inconsistencies and slang.
- In Python, two popular NLP packages are `nltk` and `spacy`.
- `nltk` is more popular but not as advanced and well maintained. `spacy` is more modern but not available for commercial use.

NATURAL LANGUAGE PROCESSING WITH 'SPACY'


- We'll be using spacy to process some news article titles. First load the NLP toolkit by specifying the language.

```
from spacy.en import English  
nlp_toolkit = English()
```

- This toolkit has 3 pre-processing engines:
 - A tokenizer: to identify the word tokens
 - A tagger: to identify the concepts described by the words
 - A parser: to identify the phrases and links between different words
- Each of these engines can be overridden with a different, specialized tool. You can even write your own and use them in place of the defaults.

NATURAL LANGUAGE PROCESSING WITH ‘SPACY’

- ▶ The first title is “[IBM Sees Holographic Calls, Air Breathing Batteries](#)”.
- ▶ From this, we may want to extract several pieces of information: this title references a company and that company is referencing a new possible product: air-breathing batteries.



NATURAL LANGUAGE PROCESSING WITH 'SPACY'

- We can use spacy to get information about this title.

```
title = "IBM sees holographic calls, air breathing batteries"
parsed = nlp_toolkit(title)
```

```
for (i, word) in enumerate(parsed):
    print("Word: {}".format(word))
    print("\t Phrase type: {}".format(word.dep_))
    print("\t Is the word a known entity type? {}".format(word.ent_type_) if
word.ent_type_ else "No"))
    print("\t Lemma: {}".format(word.lemma_))
    print("\t Parent of this word: {}".format(word.head.lemma_))
```

- nlp_toolkit runs each of the individual pre-processing tools.

NATURAL LANGUAGE PROCESSING WITH 'SPACY'

- The output will look similar to this:

Word: IBM

Phrase type: nsubj

Is the word a known entity type? ORG

Lemma: ibm

Parent of this word: see

Word: sees

Phrase type: ROOT

Is the word a known entity type? No

Lemma: see

Parent of this word: see

Word: holographic

Phrase type: amod

Is the word a known entity type? No

Lemma: holographic

Parent of this word: call

NATURAL LANGUAGE PROCESSING WITH ‘SPACY’

- In this output:
 - “IBM” is identified as an organization (ORG).
 - We identify a phrase: “holographic calls”.
 - We identify a compound noun phrase: “air breathing batteries”.
 - We identify that “see” is at the root as an action “IBM” is taking.
 - We can see that “batteries” was lemmatized to “battery”.

NATURAL LANGUAGE PROCESSING WITH 'SPACY'

- We can use this output to find all titles that discuss an organization.

```
def references_organization(title):  
    parsed = nlp(title)  
    return any([word.ent_type_ == 'ORG' for word in parsed])
```

```
data['references_organization'] =  
data['title'].fillna('').map(references_organization)
```

```
data[data['references_organization']][['title']].head()
```

ACTIVITY: KNOWLEDGE CHECK



EXERCISE

COMPLETE THE FOLLOWING TASKS

1. Using the code on the previous slide, write a function to identify titles that mention an organization (ORG) and a person (PERSON).

DELIVERABLE

New function

COMMON PROBLEMS IN NLP

- These subtasks are very difficult, because language is complex and changes frequently.
- Most often, we are looking for heuristics to search through large amounts of text data. The results may not be perfect... and that's okay!
- Older techniques rely on rule-based systems. More recent techniques use flexible systems, focusing on the words used rather than the structure of the sentence.

INTRODUCTION

TEXT CLASSIFICATION

TEXT CLASSIFICATION

- Text classification is the task of predicting which category or topic a text sample is from.
- For example, we may want to identify whether an article is a sports or business story. Or whether an article has positive or negative sentiment.
- Typically, this is done by using the text as features and the label as the target output. This is referred to as *bag-of-words* classification.
- To include text as features, we usually create a *binary* feature for each word, i.e. does this piece of text contain that word?

TEXT CLASSIFICATION

- To create binary text features, we first create a vocabulary to account for all possible words in our universe.
- As we do this, we need to consider several things.
 - Does order of words matter?
 - Does punctuation matter?
 - Does upper or lower case matter?

TEXT CLASSIFICATION

▸ This table illustrates features created from the following passage.

“It’s a great advantage not to drink among hard drinking people.”

Feature	Value
it’s	1
great	1
good	0
advantage	1
not	1
think	0
drink	1
from	0
hard	1
drinking	1

Feature	Value
people	1
withhold	0
random	0
smoke	0
among	1
whenever	0
thoughtful	0
inexhaustible	0
men	0
Nick	0

DEMO

TEXT PROCESSING IN SCIKIT-LEARN

TEXT PROCESSING IN SCIKIT-LEARN

- Scikit-learn has many pre-processing utilities that simplify tasks required to convert text into features for a model.
- These can be found in the `sklearn.preprocessing.text` package.
- We will use the StumbleUpon dataset again to perform text classification. This time, we will use the text content itself to predict whether a page is ‘evergreen’ or not.
- Open the starter code notebook to follow along.

COUNTVECTORIZER

- CountVectorizer converts a collection of text into a matrix of features. Each row will be a sample (an article or piece of text) and each column will be a text feature (usually a count or binary feature per word).
- CountVectorizer takes a column of text and creates a new dataset. It generates a feature for every word in all of the pieces of text.
- **REMEMBER:** Using all of the words can be useful, but we may need to use *regularization* to avoid overfitting. Otherwise, rare words may cause the model to overfit and not generalize.

COUNTVECTORIZER

- Instantiate a new CountVectorizer.

```
from sklearn.feature_extraction.text import CountVectorizer
```

```
vectorizer = CountVectorizer(max_features = 1000,  
                             ngram_range=(1, 2),  
                             stop_words='english',  
                             binary=True)
```

COUNTVECTORIZER PARAMETERS

- There are several parameters to utilize.
- `ngram_range` - a range of word phrases to use
 - `(1, 1)` means use all single words
 - `(1, 2)` means use all contiguous pairs of word
 - `(1, 3)` means use all triples
- `stop_words='english'`
 - Stop words are non-content words (e.g. 'to', 'the', 'it', etc). They aren't helpful for prediction, so they get removed.

COUNTVECTORIZER PARAMETERS

- `max_features=1000`
 - Maximum number of words to consider (uses the first N most frequent)
- `binary=True`
 - To use a dummy column as the entry (1 or 0, as opposed to the count). This is useful if you think a word appearing 10 times is no more important than whether the word appears at all.

COUNTVECTORIZER

- Vectorizers are like other models in scikit-learn.
 - We create a vectorizer object with the parameters of our feature space.
 - We `fit` a vectorizer to learn the vocabulary.
 - We `transform` a set of text into that feature space.

COUNTVECTORIZER

- Note: there is a distinction between *fit* and *transform*.
 - We **fit** from our training set. This is part of the model building process, so we don't look at our test set.
 - We **transform** our test set using our model fit on the training set.

COUNTVECTORIZER EXAMPLE

```
titles = data['title'].fillna('')
```

```
from sklearn.feature_extraction.text import CountVectorizer
```

```
vectorizer = CountVectorizer(max_features = 1000,  
                             ngram_range=(1, 2),  
                             stop_words='english',  
                             binary=True)
```

```
# Use `fit` to learn the vocabulary of the titles vectorizer.fit(titles)
```

```
# Use `transform` to generate the sample X word matrix - one column per  
feature (word or n-grams)
```

```
X = vectorizer.transform(titles)
```

RANDOM FOREST PREDICTION MODEL

- ▶ We can now build a random forest model to predict “evergreenness”.

```
from sklearn.ensemble import RandomForestClassifier
```

```
model = RandomForestClassifier(n_estimators = 20)
```

```
# Use `fit` to learn the vocabulary of the titles vectorizer.fit(titles)
```

```
# Use `transform` to generate the sample X word matrix - one column per feature (word or n-grams)
```

```
X = vectorizer.transform(titles)
```

```
y = data['label']
```

```
from sklearn.cross_validation import cross_val_score
```

```
scores = cross_val_score(model, X, y, scoring='roc_auc')
```

```
print('CV AUC {}, Average AUC {}'.format(scores, scores.mean()))
```

TERM FREQUENCY - INVERSE DOCUMENT FREQUENCY

- An alternative *bag-of-words* approach to CountVectorizer is a Term Frequency - Inverse Document Frequency (TF-IDF) representation.
- TF-IDF uses the product of two intermediate values, the *Term Frequency* and *Inverse Document Frequency*.

TERM FREQUENCY - INVERSE DOCUMENT FREQUENCY

- *Term Frequency* is equivalent to `CountVectorizer` features, just the number of times a word appears in the document (i.e. count).
- *Document Frequency* is the percentage of documents that a particular word appears in.
- For example, “the” would be 100% while “Syria” is much lower.
- *Inverse Document Frequency* is just $1/\text{Document Frequency}$.

TERM FREQUENCY - INVERSE DOCUMENT FREQUENCY

- Combining, $\text{TF-IDF} = \text{Term Frequency} * \text{Inverse Document Frequency}$ or $\text{TF-IDF} = \text{Term Frequency} / \text{Document Frequency}$
- The intuition is that the words that have high weight are those that either appear ***frequently*** in this document or appear ***rarely*** in other documents (and are therefore unique to this document).
- This is a good alternative to using a static set of “stop” words.

```
from sklearn.feature_extraction.text import TfidfVectorizer  
vectorizer = TfidfVectorizer()
```

ACTIVITY: KNOWLEDGE CHECK

ANSWER THE FOLLOWING QUESTIONS



EXERCISE

1. What does TF-IDF stand for?
2. What does this function do and why is it useful?
3. Use `TfidfVectorizer` to create a feature representation of the StumbleUpon titles.

DELIVERABLE

Answers to the above questions and feature representation

INDEPENDENT PRACTICE

TEXT CLASSIFICATION IN SCIKIT-LEARN

ACTIVITY: TEXT CLASSIFICATION IN SCIKIT-LEARN



EXERCISE

DIRECTIONS (30 minutes)

1. Use the text features of `title` with one or more feature sets from the previous random forest model. Train this model to see if it improves AUC.
2. Use the body text instead of the `title`. Does this give an improvement?
3. Use `TfidfVectorizer` instead of `CountVectorizer`. Does this give an improvement?

Check: Were you able to prepare a model that uses both quantitative features and text features? Does this model improve the AUC?

DELIVERABLE

Three new models

CONCLUSION

TOPIC REVIEW

LET'S REVIEW

- Natural language processing (NLP) is the task of pulling meaning and information from text.
- This typically involves many subproblems including tokenization, cleaning (stemming and lemmatization), and parsing.
- After we have structured our text, we can identify features for other tasks, including classification, summarization, and translation.
- In scikit-learn, we use vectorizers to create text features for classification, such as `CountVectorizer` and `TfidfVectorizer`.